

# SQL Query Optimization

## Agenda:

In today's session, we'll cover essential topics, including:-

- ◆ Busting a common myth about the COUNT()
- ◆ Selecting only the columns that you truly need
- ◆ Why LIMIT can be a hidden pitfall
- ◆ Harnessing EXISTS() instead of COUNT()
- ◆ Leveraging APPROX\_COUNT\_DISTINCT for efficient counting
- ◆ Swapping Self-Joins with Windows Functions
- ◆ The art of trimming your data early and frequently
- ◆ When to choose MAX() over RANK()
- ◆ Importance of ordering your JOINS, from larger to smaller tables
- ◆ Does the WHERE clause sequence really matter?
- ◆ Should we push ORDER BY to the end of the query?

## Summary of Previous Lecture:

### Ad-hoc reporting

In the data analysis world, being asked questions, exploring a database, writing SQL statements to find and pull the data needed to determine the answers, and conducting the analysis of that data to calculate the answers to the questions, is called **ad-hoc reporting**.

### Common Table Expressions

- Common Table Expressions (**CTEs**) allow us to create temporary named results sets that exist temporarily within the execution scope of SQL statements such as SELECT, INSERT, UPDATE, DELETE, and MERGE.
- The **WITH** clause in SQL provides a better way to write the auxiliary/helper statement which can be later used in larger queries.
- These statements are referred to as common table expressions which are nothing but defining a temporary relational table to be used later in a SQL statement.
- The table is being called temporary because it exists only during the scope of the SQL statement written after CTEs.

**The syntax for CTEs is:**

```
WITH [query_alias] AS (
    [query]
),
[query_2_alias] AS (
    [query_2]
)
SELECT [column list]
FROM [query_alias]

... [remainder of query that references aliases created above]
```

-> where “[query\_alias]” is a placeholder for the name you want to use to refer to a query later, and “[query]” is a placeholder for the query you want to reuse.

### **Advantages of CTE:**

- Making recursive queries.
- Hold a query output virtually in a temporary area named as given while definition.
- No need to save Metadata.
- Useful when there is a need to do more operations on some query output.

## **Views**

- Another approach to CTEs is Views.
- A view is treated just like a table in SQL, the only difference being that it has run when it's referenced to dynamically generate a result set (where a table stores the data instead of storing the query),
- So queries that reference views can take longer to run than queries that reference tables.
- However, the view is retrieving the latest data from the underlying tables each time it is run, so you are working with the freshest data available when you query from a view.

### **Syntax:**

```
CREATE VIEW [db_name.]view_name [(column_list)]
AS
    select-statement;
• If you want to store your dataset as a view, you simply precede your SELECT statement by replacing the bracketed statements with the actual schema name, and the name you are giving the view.
```

### **When should you use each one?**

Although some differences exist, common table expressions and views seem to perform very similarly.

- **Ad-hoc queries.**
  - For queries referenced occasionally (or just once), it's usually better to use a CTE.
  - If you need the query again, you can just copy the CTE and modify it if necessary.
- **Frequently used queries.**
  - Creating a corresponding view is a good idea if you often reference the same query.
  - However, you'll need to create view permission in your database to create a view.
- **Access management.**
  - A view might restrict particular users' database access while allowing them to get the necessary information.
  - You can give users access to specific views that query the data they're allowed to see without exposing the whole database.
  - In such a case, a view provides an additional access layer.